


```
from google.colab import files
uploaded=files.upload()
```

 Choose Files pima.csv

- **pima.csv**(text/csv) - 23279 bytes, last modified: 7/8/2025 - 100% done  
Saving pima.csv to pima (1).csv

```
import pandas
import numpy
import matplotlib.pyplot as plt
from pandas.plotting import scatter_matrix
from sklearn import model_selection
from sklearn.preprocessing import StandardScaler
```

```
from sklearn.metrics import classification_report
from sklearn.metrics import confusion_matrix
from sklearn.metrics import accuracy_score
```


```
from sklearn.pipeline import Pipeline
from sklearn.model_selection import GridSearchCV
```

```
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
from sklearn.tree import DecisionTreeClassifier
from sklearn.svm import SVC
```

```
from sklearn.ensemble import AdaBoostClassifier
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.ensemble import ExtraTreesClassifier
```

```
url="pima.csv"
names=['preg','plas','pres','skin','test','mass','pedi','age','class']
data=pandas.read_csv(url,names=names)
```


```
print(data)
```



	preg	plas	pres	skin	test	mass	pedi	age	class
0	6	148	72	35	0	33.6	0.627	50	1
1	1	85	66	29	0	26.6	0.351	31	0
2	8	183	64	0	0	23.3	0.672	32	1
3	1	89	66	23	94	28.1	0.167	21	0
4	0	137	40	35	168	43.1	2.288	33	1
..	...	...	...	...	...	...	...	...	...
763	10	101	76	48	180	32.9	0.171	63	0
764	2	122	70	27	0	36.8	0.340	27	0
765	5	121	72	23	112	26.2	0.245	30	0
766	1	126	60	0	0	30.1	0.349	47	1
767	1	93	70	31	0	30.4	0.315	23	0

```
[768 rows x 9 columns]
```

```
print(data.dtypes)
print(data.shape)
```



```
preg      int64
plas      int64
pres      int64
skin      int64
test      int64
mass      float64
pedi      float64
age       int64
class     int64
dtype: object
(768, 9)
```

```
#checking for null values
data.apply(lambda x:sum(x.isnull()),axis=0)
```

```

0
preg    0
plas    0
pres    0
skin    0
test    0
mass    0
pedi    0
age     0
class   0

dtype: int64
```

```
pandas.set_option('display.precision',2)
print(data.describe())
```

```


count    preg    plas    pres    skin    test    mass    pedi    age    class
mean      3.85  120.89   69.11   20.54   79.80   31.99   0.47   33.24   0.35
std       3.37   31.97   19.36   15.95  115.24    7.88   0.33   11.76   0.48
min       0.00    0.00    0.00    0.00    0.00    0.00   0.08   21.00    0.00
25%       1.00   99.00   62.00    0.00    0.00   27.30   0.24   24.00    0.00
50%       3.00  117.00   72.00   23.00   30.50   32.00   0.37   29.00    0.00
75%       6.00  140.25   80.00   32.00  127.25   36.60   0.63   41.00    1.00
max      17.00  199.00  122.00   99.00  846.00   67.10   2.42   81.00    1.00
```

```
print("class size",data.groupby('class').size())
```

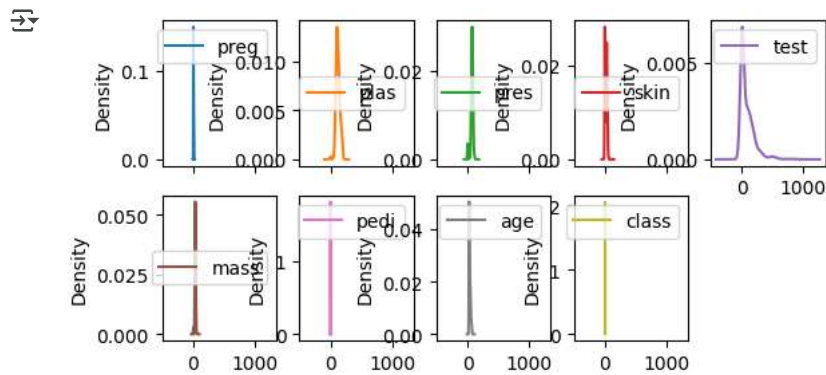
```

class size class
0      500
1      268
dtype: int64
```


```
data.hist()
plt.show()
```




```
data.plot(kind='density',subplots=True,layout=(3,5))
plt.show()
```




```
data.plot(kind='box',subplots=True,layout=(3,3))
plt.show()
```



```
scatter_matrix(data)
plt.show()
```



```
fig=plt.figure()
ax=fig.add_subplot(111)
cax=ax.matshow(data.corr(),vmin=-1,vmax=1)
fig.colorbar(cax)
plt.show()
```



```
array=data.values
X=array[:,0:8]
Y=array[:,8]
```

```
models=[]
models.append(('LR',LogisticRegression()))
models.append(('LDA',LinearDiscriminantAnalysis()))
models.append(('SVM',SVC()))
models.append(('CART',DecisionTreeClassifier()))
models.append(('KNN',KNeighborsClassifier()))
models.append(('NB',GaussianNB()))
results=[]
names=[]
scoring='accuracy'
test_size=0.33
seed=8
for name,model in models:
    kfold=model_selection.ShuffleSplit(n_splits=10,test_size=test_size,random_state=seed)
    cv_results=model_selection.cross_val_score(model,X, Y,cv=kfold,scoring=scoring)
    results.append(cv_results)
    names.append(name)
    msg="%s:%f(%f)"%(name,cv_results.mean()*100,cv_results.std())
    print(msg)
```



Please also refer to the documentation for alternative solver options:  
<https://scikit-learn.org/stable/modules/preprocessing.html>  
[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

```
/usr/local/lib/python3.11/dist-packages/sklearn/linear_model/_logistic.py:465: ConvergenceWarning: lbfgs failed to converge (status=1)
STOP: TOTAL NO. OF ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

```
/usr/local/lib/python3.11/dist-packages/sklearn/linear_model/_logistic.py:465: ConvergenceWarning: lbfgs failed to converge (status=1)
STOP: TOTAL NO. OF ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

```
LR:76.456693(0.021906)
```

```
LDA:76.614173(0.018144)
```

```
SVM:74.448819(0.018836)
```

```
CART:67.401575(0.029399)
```


```
KNN:71.850394(0.019307)
```

```
NR:75.039370(0.019935)
```

```
fig=plt.figure()
fig.suptitle('Algorithm Comparison')
ax=fig.add_subplot(111)
plt.boxplot(results)
ax.set_xticklabels(names)
plt.show()
```



Algorithm Comparison




```
pipelines=[]
pipelines.append(('ScaledLR',Pipeline([('Scaler',StandardScaler()),('LR',LogisticRegression())]))
pipelines.append(('ScaledLDA',Pipeline([('Scaler',StandardScaler()),('LDA',LinearDiscriminantAnalysis())]))
pipelines.append(('ScaledSVM',Pipeline([('Scaler',StandardScaler()),('SVM',SVC())]))
pipelines.append(('ScaledKNN',Pipeline([('Scaler',StandardScaler()),('KNN',KNeighborsClassifier())]))
pipelines.append(('ScaledCART',Pipeline([('Scaler',StandardScaler()),('CART',DecisionTreeClassifier())]))
pipelines.append(('ScaledNB',Pipeline([('Scaler',StandardScaler()),('NB',GaussianNB())]))
results=[]
names=[]
scoring='accuracy'
test_size=0.33
seed=8
for name,model in pipelines:
    kfold=model_selection.ShuffleSplit(n_splits=10,test_size=test_size,random_state=seed)
    cv_results=model_selection.cross_val_score(model,X,Y,cv=kfold,scoring=scoring)
    results.append(cv_results)
    names.append(name)
```

```
msg="%s:%f(%f)"%(name,cv_results.mean()*100,cv_results.std())
print(msg)
```

```
→ ScaledLR:76.692913(0.019110)
ScaledLDA:76.614173(0.018144)
ScaledSVM:75.472441(0.022205)
ScaledKNN:71.968504(0.020518)
ScaledCART:68.740157(0.024850)
ScaledNB:75.039370(0.019935)
```

```
fig=plt.figure()
fig.suptitle('Scaled Algorithm Comparision')
ax=fig.add_subplot(111)
plt.boxplot(results)
ax.set_xticklabels(names)
plt.show()
```

→ **Scaled Algorithm Comparision**



```
Scaler=StandardScaler().fit(X)
rescaledX=Scaler.transform(X)
neighbors=[1,3,5,7,9,11,13,15,17,19,21]
model=KNeighborsClassifier()
param_grid=dict(n_neighbors=neighbors)
kfold=model_selection.ShuffleSplit(n_splits=10,test_size=test_size,random_state=seed)
grid=GridSearchCV(estimator=model,param_grid=param_grid,scoring=scoring,cv=kfold)
grid_result=grid.fit(rescaledX,Y)
print("Best:%f using %s"%(grid_result.best_score_,grid_result.best_params_))
```

→ Best:0.743701 using {'n\_neighbors': 17}

```
Scaler=StandardScaler().fit(X)
rescaledX=Scaler.transform(X)
c_values=[0.1,0.3,0.5,0.7,0.9,1.1,1.3,1.5,1.7,3.0]
kernel_values=['linear','poly','rbf','sigmoid']
param_grid=dict(C=c_values,kernel=kernel_values)
model=SVC()
kfold=model_selection.ShuffleSplit(n_splits=10,test_size=test_size,random_state=seed)
grid=GridSearchCV(estimator=model,param_grid=param_grid,scoring=scoring,cv=kfold)
grid_result=grid.fit(rescaledX,Y)
print("Best%f using %s"%(grid_result.best_score_,grid_result.best_params_))
```

→ Best0.768110 using {'C': 0.3, 'kernel': 'linear'}:


```
ensembles=[]
ensembles.append(('AB',AdaBoostClassifier()))
ensembles.append(('GBM',GradientBoostingClassifier()))
ensembles.append(('RF',RandomForestClassifier()))
ensembles.append(('ET',ExtraTreesClassifier()))
results=[]
names=[]
scoring='accuracy'
```

```
test_size=0.33
seed=8
for name,model in ensembles:
    kfold=model_selection.ShuffleSplit(n_splits=10,test_size=test_size,random_state=seed)
    cv_results=model_selection.cross_val_score(model,X, Y,cv=kfold,scoring=scoring)
    results.append(cv_results)
    names.append(name)
    msg="%s:%f(%f)"%(name,cv_results.mean()*100,cv_results.std())
    print(msg)
```

```
AB:74.842520(0.025481)
GBM:75.000000(0.013892)
RF:75.039370(0.016156)
ET:74.803150(0.021420)
```

```
fig=plt.figure()
fig.suptitle("Ensemble Algorithm Comparision")
ax=fig.add_subplot(111)
plt.boxplot(results)
ax.set_xticklabels(names)
plt.show()
```

Ensemble Algorithm Comparision



```
Scaler=StandardScaler().fit(X)
rescaledX=Scaler.transform(X)
model=LogisticRegression()
model.fit(rescaledX,Y)
```

```
LogisticRegression
```

```
rescaledValidationX=Scaler.transform(X)
predictions=model.predict(rescaledValidationX)
```

```
print(classification_report(Y,predictions))
```

```
precision    recall  f1-score   support

0.0         0.80      0.89      0.84        500
1.0         0.74      0.58      0.65        268

accuracy          0.78          768
macro avg         0.77         0.74         0.75          768
weighted avg         0.78         0.78         0.78          768
```

```
print(accuracy_score(Y,predictions))
```

```
0.7838541666666666
```

```
print(confusion_matrix(Y,predictions))
```

```
↵ [[446  54]  
   [112 156]]
```